

Neural Networks and Deep Learning

ICP7

Student Name: Sravani Lankala
Student Id: 700746285

GitHub Link: https://github.com/sravanilankala/NNDL_ICP7_Fall2023

Video Link:

https://drive.google.com/file/d/1Ifp6Cq_4ZWkUUlkowsTI5Ue8JH90S4wK/view?usp=sharing

1. Follow the instruction below and then report how the performance changed.(apply all at once)
 - Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2.
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2.
 - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer,128 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2.
 - Flatten layer.
 - Dropout layer at 20%.
 - Fully connected layer with 1024 units and a rectifier activation function.
 - Dropout layer at 20%.
 - Fully connected layer with 512 units and a rectifier activation function.
 - Dropout layer at 20%.
 - Fully connected output layer with 10 units and a Softmax activation function

+ Code + Text

```

✓ 1m # 1. Follow the instruction below and then report how the performance changed.(apply all at once)
#• Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
#• Dropout layer at 20%.
#• Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
#• Max Pool layer with size 2×2.
#• Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
#• Dropout layer at 20%.
#• Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
#• Max Pool layer with size 2×2.
#• Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
#• Dropout layer at 20%.
#• Convolutional layer,128 feature maps with a size of 3×3 and a rectifier activation function.
#• Max Pool layer with size 2×2.
#• Flatten layer.
#• Dropout layer at 20%.
#• Fully connected layer with 1024 units and a rectifier activation function.
#• Dropout layer at 20%.
#• Fully connected layer with 512 units and a rectifier activation function.
#• Dropout layer at 20%.
#• Fully connected output layer with 10 units and a Softmax activation function

import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.constraints import MaxNorm as maxnorm
from keras import utils as np_utils
    
```

+ Code + Text

```

✓ 1m from keras.optimizers.legacy import SGD

# Fix random seed for reproducibility
np.random.seed(7)

# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()


# Normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# One hot encode outputs
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu'))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
    
```

+ Code + Text

```

1m  model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 5
learning_rate = 0.01
decay_rate = learning_rate / epochs
sgd = SGD(lr=learning_rate, momentum=0.9, decay=decay_rate, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())


# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))

```

+ Code + Text

```

1m  conv2d_12 (Conv2D)      (None, 32, 32, 32)      896
dropout_12 (Dropout)    (None, 32, 32, 32)      0
conv2d_13 (Conv2D)      (None, 32, 32, 32)      9248
max_pooling2d_6 (MaxPoolin (None, 16, 16, 32)      0
g2D)
conv2d_14 (Conv2D)      (None, 16, 16, 64)      18496
dropout_13 (Dropout)    (None, 16, 16, 64)      0
conv2d_15 (Conv2D)      (None, 16, 16, 64)      36928
max_pooling2d_7 (MaxPoolin (None, 8, 8, 64)        0
g2D)
conv2d_16 (Conv2D)      (None, 8, 8, 128)       73856
dropout_14 (Dropout)    (None, 8, 8, 128)       0
conv2d_17 (Conv2D)      (None, 8, 8, 128)       147584
max_pooling2d_8 (MaxPoolin (None, 4, 4, 128)       0
g2D)
flatten_2 (Flatten)     (None, 2048)            0
dropout_15 (Dropout)    (None, 2048)            0

```

NNDL_700746285_ICP7.ipynb ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```

1m
dense_6 (Dense) (None, 1024) 2098176
dropout_16 (Dropout) (None, 1024) 0
dense_7 (Dense) (None, 512) 524800
dropout_17 (Dropout) (None, 512) 0
dense_8 (Dense) (None, 10) 5130

=====
Total params: 2915114 (11.12 MB)
Trainable params: 2915114 (11.12 MB)
Non-trainable params: 0 (0.00 Byte)

None
Epoch 1/5
1563/1563 [=====] - 14s 8ms/step - loss: 1.9144 - accuracy: 0.2887 - val_loss: 1.6637 - val_accuracy: 0.3998
Epoch 2/5
1563/1563 [=====] - 12s 7ms/step - loss: 1.5480 - accuracy: 0.4342 - val_loss: 1.4825 - val_accuracy: 0.4573
Epoch 3/5
1563/1563 [=====] - 11s 7ms/step - loss: 1.4103 - accuracy: 0.4871 - val_loss: 1.3557 - val_accuracy: 0.5014
Epoch 4/5
1563/1563 [=====] - 11s 7ms/step - loss: 1.3326 - accuracy: 0.5167 - val_loss: 1.2915 - val_accuracy: 0.5369
Epoch 5/5
1563/1563 [=====] - 12s 8ms/step - loss: 1.2704 - accuracy: 0.5406 - val_loss: 1.2491 - val_accuracy: 0.5455
Accuracy: 54.55%

```

Did the performance change?

With the addition of more layers and feature maps, the model's performance is likely to improve, but the complexity and training time of the model will also rise. The new model architecture described in the instruction has more feature maps and new layers, which could increase the model's accuracy.

- Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly.

NNDL_700746285_ICP7.ipynb ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```

0s
# 2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4
# images to check whether or not the model has predicted correctly.

# Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])

# Convert the predictions to class labels
predicted_labels = np.argmax(predictions, axis=1)

# Convert the actual labels to class labels
actual_labels = np.argmax(y_test[:4], axis=1)

# Print the predicted and actual labels for the first 4 images
print("Predicted labels:", predicted_labels)
print("Actual labels:", actual_labels)

1/1 [=====] - 0s 107ms/step
Predicted labels: [3 1 8 8]
Actual labels: [3 8 8 0]

```

- Visualize Loss and Accuracy using the history object

+ Code + Text

✓
1s



3. Visualize Loss and Accuracy using the history object

```
import matplotlib.pyplot as plt
```

```
# Plot the training and validation loss
```

```
plt.figure(figsize=(12, 6))
```

```
plt.subplot(1, 2, 1)
```

```
plt.plot(history.history['loss'], label='Training Loss')
```

```
plt.plot(history.history['val_loss'], label='Validation Loss')
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Loss')
```

```
plt.title('Training and Validation Loss')
```

```
plt.legend()
```

```
# Plot the training and validation accuracy
```

```
plt.subplot(1, 2, 2)
```

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
```

```
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Accuracy')
```

```
plt.title('Training and Validation Accuracy')
```

```
plt.legend()
```

```
plt.tight_layout()
```

```
plt.show()
```

+ Code + Text

✓
1s

